

23) Decision Trees

Vitor Kamada

March 2018

Tables, Graphics, and Figures from
An Introduction to Statistical Learning

James et al. (2017): Chapters: 8.1, 8.3.1, and
8.3.2

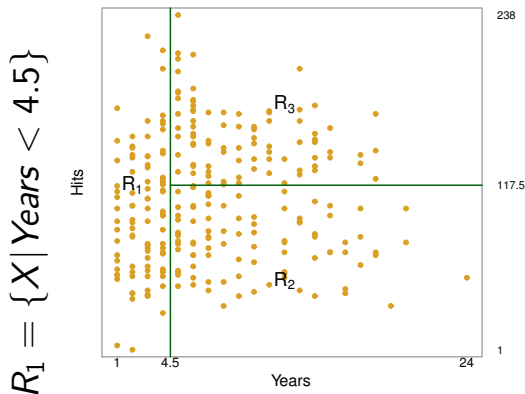
Hitters Data



$$e^{5.11} \cong \$165K, e^6 \cong \$402K, e^{6.74} \cong \$845K$$

Three-Region Partition

$$R_3 = \{X \mid \text{Years} \geq 4.5, \text{Hits} \geq 117.5\}$$



$$R_2 = \{X \mid \text{Years} \geq 4.5, \text{Hits} < 117.5\}$$

Top-down Greedy (Recursive Binary Splitting)

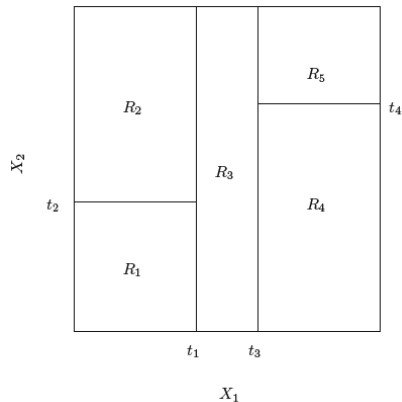
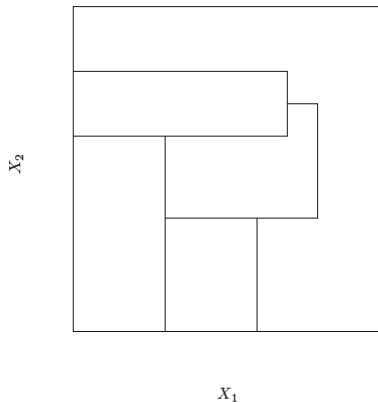
$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

$$R_1(j, s) = \{X | X_j < s\}$$

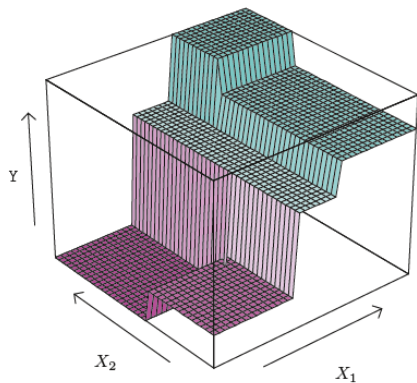
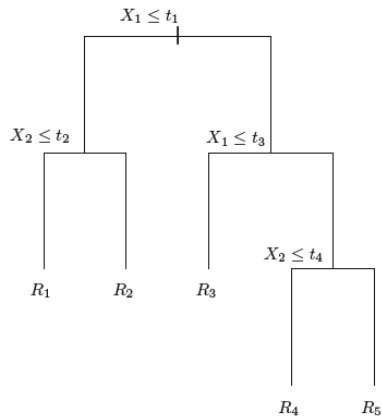
$$R_2(j, s) = \{X | X_j \geq s\}$$

$$\sum_{i: X_i \in R_1(j, s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i: X_i \in R_2(j, s)} (y_i - \hat{y}_{R_2})^2$$

No Recursive Binary Splitting vs Recursive Binary Splitting



Tree and Perspective Plot

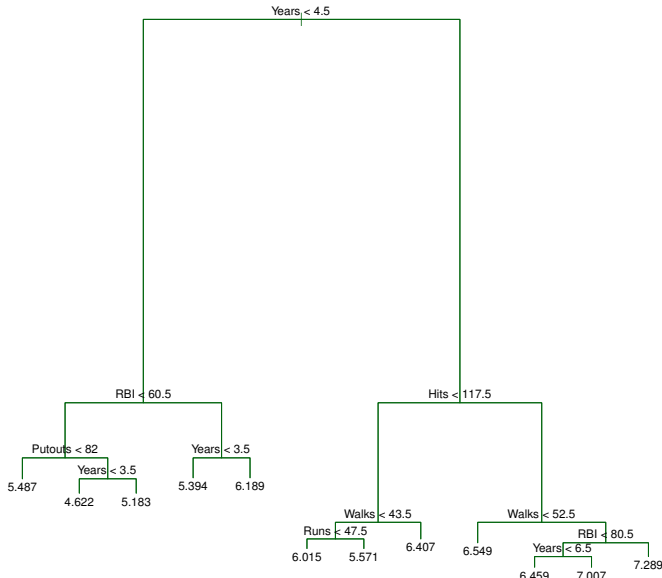


Cost Complexity Pruning (Weakest Link Pruning)

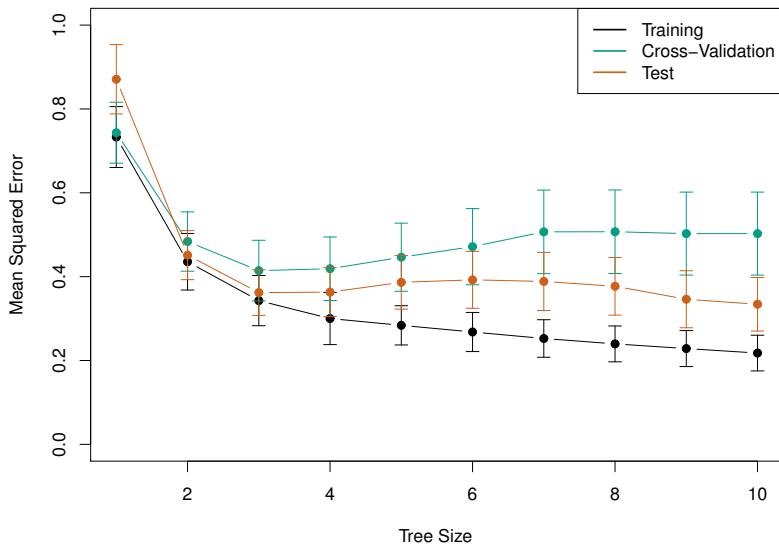
$$\sum_{m=1}^{|T|} \sum_{x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T|$$

$|T| = \#$ of terminal nodes of the tree

Unpruned Tree (Top-down Greedy Splitting)



Six-Fold Cross-Validation for Pruning Tree



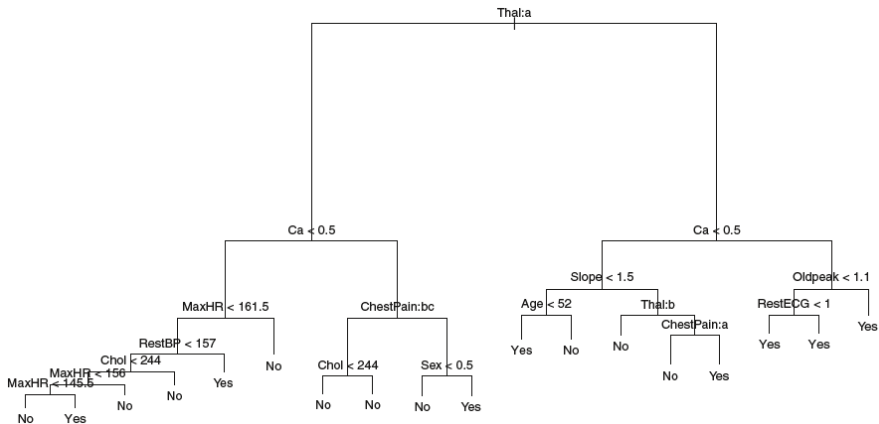
Training Error Rate, Gini Index, and Entropy

$$E = 1 - \max_k (\hat{p}_{mk})$$

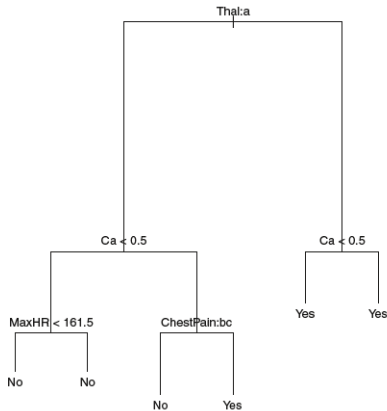
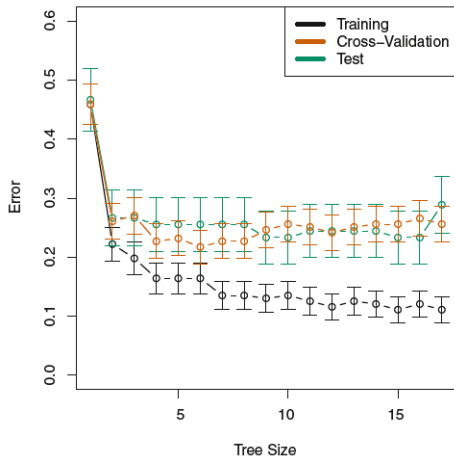
$$G = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$

$$D = - \sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}$$

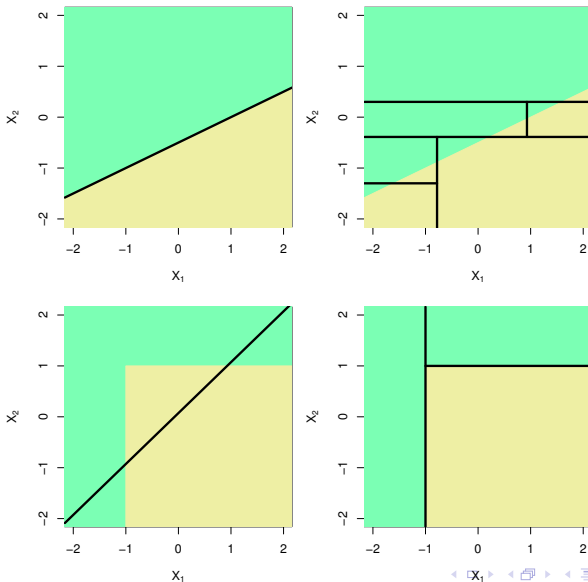
Heart Data: Unpruned Tree



Pruned Tree (Minimal Cross-Validation Error)



Linear vs Non-linear True Decision Boundary



library(ISLR); library(stargazer); stargazer(Carseats)

Statistic	N	Mean	St. Dev.	Min	Max
Sales	400	7.496	2.824	0.000	16.270
CompPrice	400	124.975	15.335	77	175
Income	400	68.657	27.986	21	120
Advertising	400	6.635	6.650	0	29
Population	400	264.840	147.376	10	509
Price	400	115.795	23.677	24	191
Age	400	53.322	16.200	25	80
Education	400	13.900	2.621	10	18

```
library(tree); High=ifelse(Sales<=8,"No","Yes")
```

```
Carseats=data.frame(Carseats,High)
```

```
tree.carseats=tree(High~.-Sales,Carseats);
```

```
summary(tree.carseats)
```

Classification tree:

```
tree(formula = High ~ . - Sales, data = Carseats)
```

Variables actually used in tree construction:

```
[1] "ShelveLoc" "Price" "Income" "CompPrice"
```

```
[5] "Population" "Advertising" "Age" "US"
```

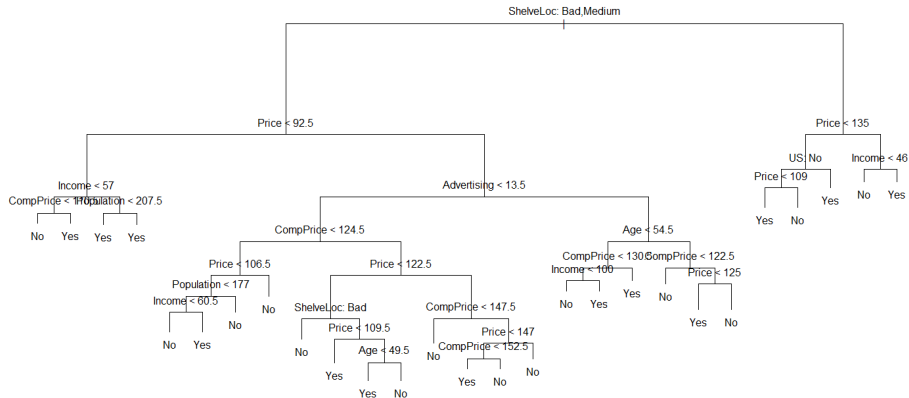
Number of terminal nodes: 27

Residual mean deviance: 0.4575 = 170.7 / 373

Misclassification error rate: 0.09 = 36 / 400

$$\frac{-2 \sum_k \sum_m n_{mk} \log \hat{p}_{mk}}{n - |T_0|} = \frac{170.7}{400 - 27} = 0.45$$

plot(tree.carseats); text(tree.carseats,pretty=0)



node), split, n, deviance, yval, (yprob)

* denotes terminal node

- 1) root 400 541.500 No (0.59000 0.41000)
 - 2) ShelfLoc: Bad,Medium 315 390.600 No (0.68889 0.31111)
 - 4) Price < 92.5 46 56.530 Yes (0.30435 0.69565)
 - 8) Income < 57 10 12.220 No (0.70000 0.30000)
 - 16) CompPrice < 110.5 5 0.000 No (1.00000 0.00000) *
 - 17) CompPrice > 110.5 5 6.730 Yes (0.40000 0.60000) *
 - 9) Income > 57 36 35.470 Yes (0.19444 0.80556)
 - 18) Population < 207.5 16 21.170 Yes (0.37500 0.62500) *
 - 19) Population > 207.5 20 7.941 Yes (0.05000 0.95000) *
 - 5) Price > 92.5 269 299.800 No (0.75465 0.24535)
 - 10) Advertising < 13.5 224 213.200 No (0.81696 0.18304)
 - 20) CompPrice < 124.5 96 44.890 No (0.93750 0.06250)
 - 40) Price < 106.5 38 33.150 No (0.84211 0.15789)
 - 80) Population < 177 12 16.300 No (0.58333 0.41667)
 - 160) Income < 60.5 6 0.000 No (1.00000 0.00000) *
 - 161) Income > 60.5 6 5.407 Yes (0.16667 0.83333) *
 - 81) Population > 177 26 8.477 No (0.96154 0.03846) *
 - 41) Price > 106.5 58 0.000 No (1.00000 0.00000) *

```
set.seed(2); train=sample(1:nrow(Carseats), 200)
```

```
Carseats.test=Carseats[-train,]
```

```
High.test=High[-train]
```

```
tree.carseats=tree(High~.-Sales,Carseats,subset=train)
```

```
tree.pred=predict(tree.carseats,Carseats.test,type="class")
```

```
table(tree.pred,High.test)
```

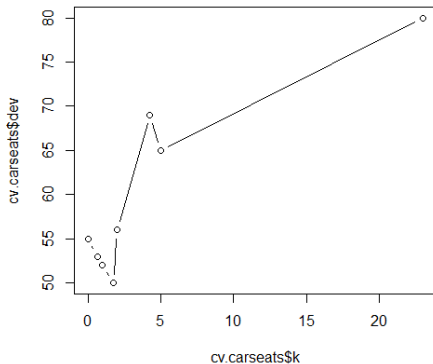
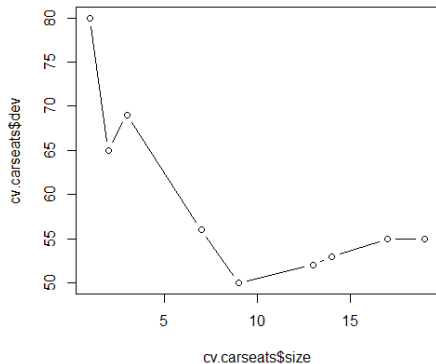
	High.test	
tree.pred	No	Yes
No	86	27
Yes	30	57

$$\frac{86+57}{200} = 71.5\%$$

```
set.seed(3); cv.carseats=cv.tree(tree.carseats,  
FUN=prune.misclass); par(mfrow=c(1,2))
```

```
plot(cv.carseats$size,cv.carseats$dev,type="b")
```

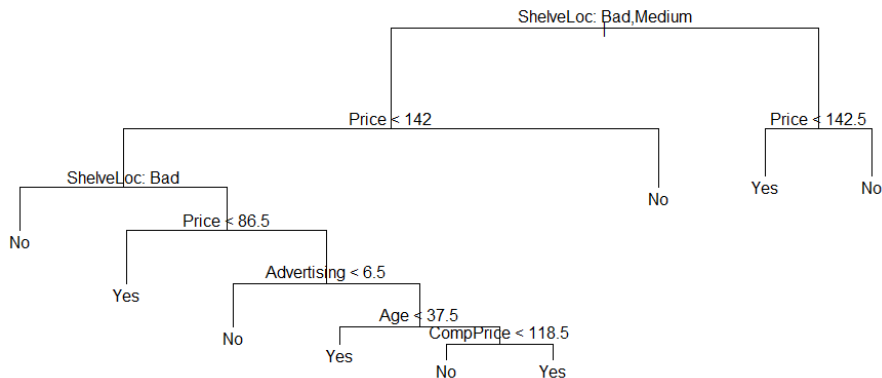
```
plot(cv.carseats$k,cv.carseats$dev,type="b")
```



```
prune.carseats=prune.misclass(tree.carseats,best=9)
```

```
plot(prune.carseats)
```

```
text(prune.carseats,pretty=0)
```



```
tree.pred=predict(prune.carseats,  
Carseats.test,type="class")
```

```
table(tree.pred,High.test)
```

	High.test	
tree.pred	No	Yes
No	94	24
Yes	22	60

$$\frac{94+60}{200} = 77\%$$

```
prune.carseats=prune.misclass(tree.carseats,best=15)
```

```
tree.pred=predict(prune.carseats,  
Carseats.test,type="class")
```

```
table(tree.pred,High.test)
```

	High.test	
tree.pred	No	Yes
No	86	22
Yes	30	62

$$\frac{86+62}{200} = 74\%$$