

25) Support Vector Machines (SVM)

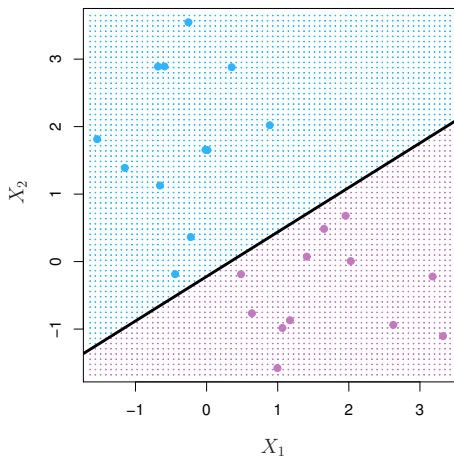
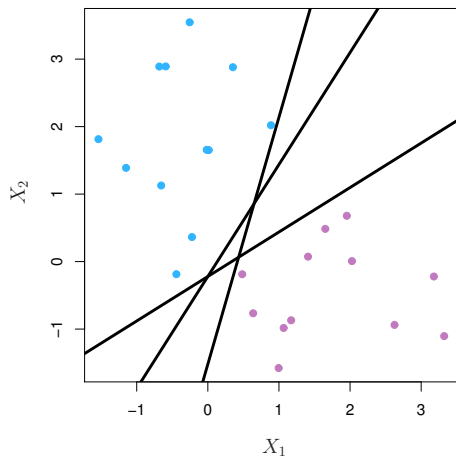
Vitor Kamada

April 2018

Tables, Graphics, and Figures from
An Introduction to Statistical Learning

James et al. (2017): Chapters: 9.3, 9.4, 9.6.2,
9.6.3, 9.6.4

Separating Hyperplane



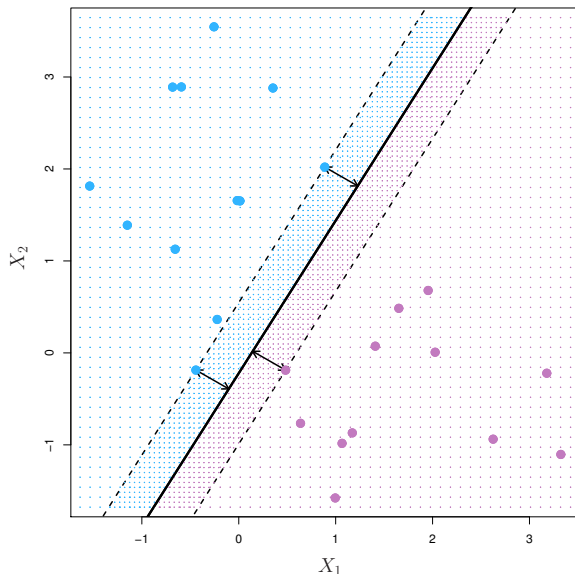
Hyperplane

$$\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip} > 0 \text{ if } y_i = 1$$

$$\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip} < 0 \text{ if } y_i = -1$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) > 0$$

Maximal Margin Classifier




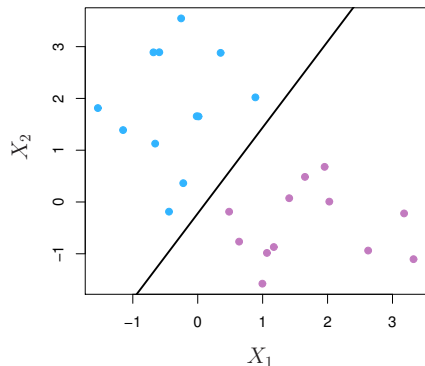
Construction of the Maximal Margin Classifier

$$\underset{\beta_0, \beta_1, \dots, \beta_p, M}{\text{maximize}} \quad M$$

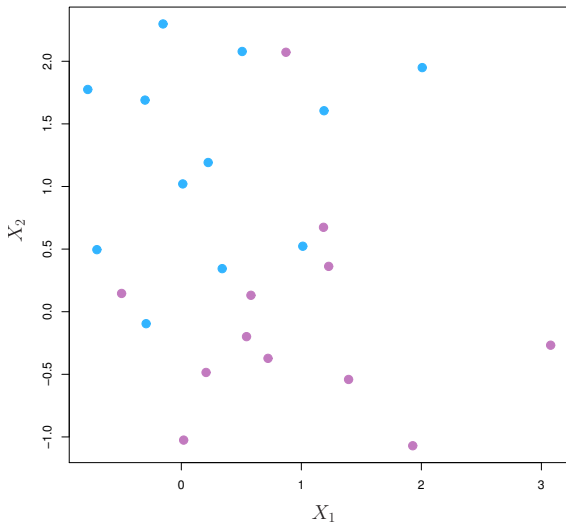
$$\text{subject to} \quad \sum_{j=1}^p \beta_j^2 = 1$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M$$

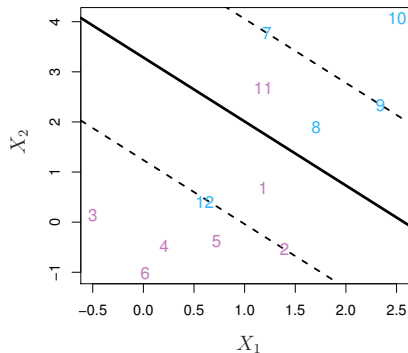
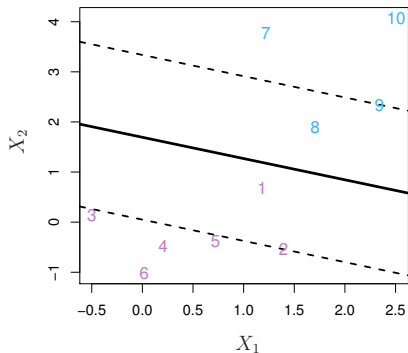
Maximal Margin Hyperplane



No Separating Hyperplane



Soft Margin Classifier



Construction of Support Vector Classifier

$$\underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n, M}{\text{maximize}} \quad M$$

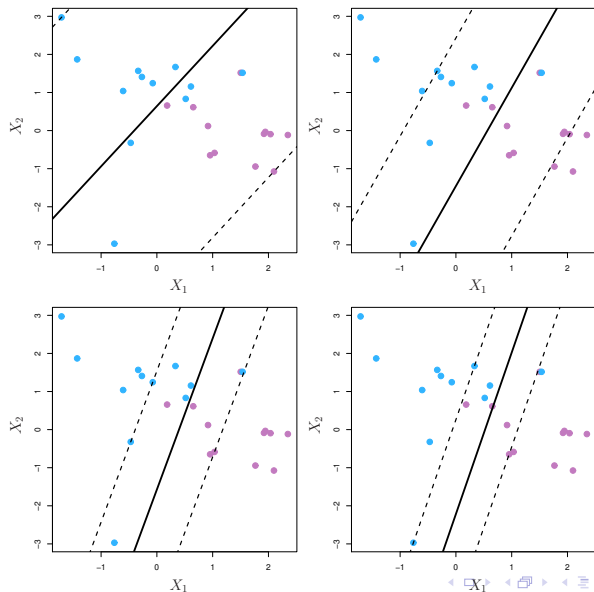
$$\text{subject to } \sum_{j=1}^p \beta_j^2 = 1$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i)$$


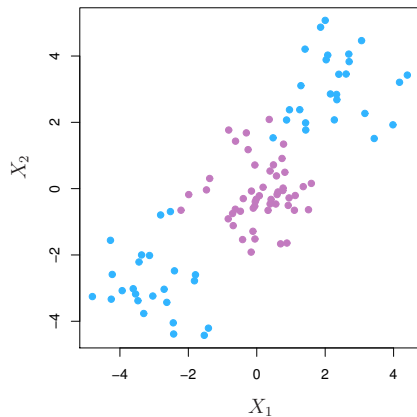
$$\epsilon_i \geq 0$$

$$\sum_{i=1}^n \epsilon_i \leq C$$

Different Values of the Tuning Parameter C



Non-Linear Boundary



$$\underset{\beta_0, \beta_{11}, \beta_{12}, \dots, \beta_{p1}, \beta_{p2}, \epsilon_1, \dots, \epsilon_n, M}{\text{maximize}}$$

$$\text{subject to } \sum_{j=1}^p \sum_{k=1}^2 \beta_{jk}^2 = 1$$

$$y_i \left(\beta_0 + \sum_{j=1}^p \beta_{j1} x_{ij} + \sum_{j=1}^p \beta_{j2} x_{ij}^2 \right) \geq M(1 - \epsilon_i)$$

$$\epsilon_i \geq 0$$

$$\sum_{i=1}^n \epsilon_i \leq C$$



$$\text{Kernel} = K(x_i, x_{i'})$$

Linear: $\sum_{j=1}^p (x_{ij}, x_{i'j})$



Polynomial: $(1 + \sum_{j=1}^p x_{ij}, x_{i'j})^d$

Radial: $\exp[-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2]$



Polynomial Kernel of Degree 3 vs Radial Kernel



Heart Data Training Set



Heart Data Test Set

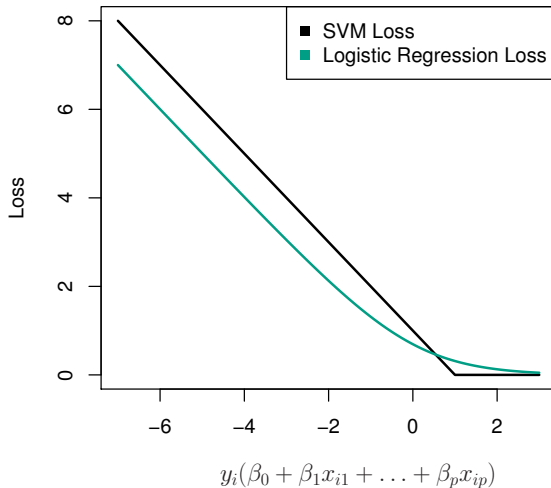


Support Vector Machine Loss Function

$$f(X) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$$

$$\underset{\beta_0, \beta_1, \dots, \beta_p}{\text{minimize}} \left\{ \sum_{i=1}^n \max[0, 1 - y_i f(x_i)] + \lambda \sum_{j=1}^p \beta_j^2 \right\}$$

SVM vs Logistic Regression Loss Function



```
library(e1071); set.seed(1)
```

```
x=matrix(rnorm(200*2), ncol=2)
```

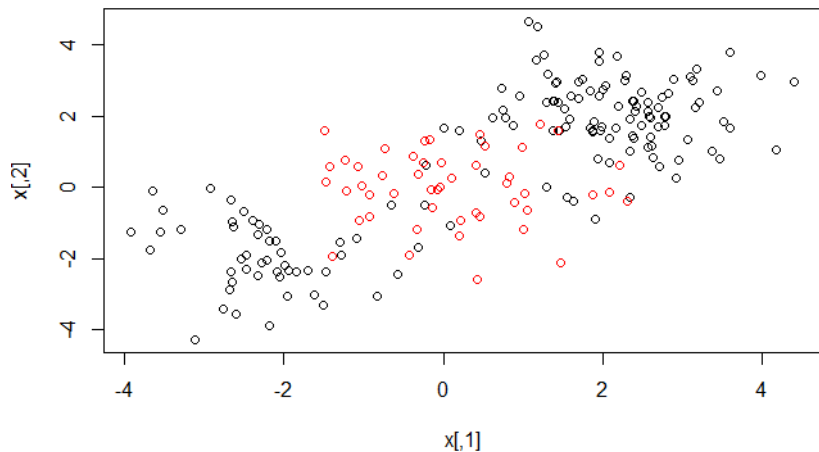
```
x[1:100,]=x[1:100,]+2
```

```
x[101:150,]=x[101:150,]-2
```

```
y=c(rep(1,150),rep(2,50))
```

```
dat=data.frame(x=x,y=as.factor(y))
```


`plot(x, col=y)`



```
train=sample(200,100)
```


```
svmfit=svm(y~., data=dat[train,], kernel="radial",  
gamma=1, cost=1); plot(svmfit, dat[train,])
```

SVM classification plot



```
svmfit=svm(y~., data=dat[train,],  
kernel="radial",gamma=1,cost=1e5)
```

SVM classification plot



```
tune.out=tune(svm, y~., data=dat[train,], kernel="radial",
ranges=list(cost=c(0.1,1,10,100,1000),gamma=c(0.5,1,2,3,4)))
```

summary(tune.out)

```
- sampling method: 10-fold cross validation

- best parameters:
  cost gamma
    1    0.5
10 1e+03    1.0  0.14 0.12649111
11 1e-01    2.0  0.18 0.12292726
12 1e+00    2.0  0.10 0.08164966
13 1e+01    2.0  0.12 0.09189366
14 1e+02    2.0  0.19 0.12866839
15 1e+03    2.0  0.18 0.13165612
16 1e-01    3.0  0.22 0.13165612
17 1e+00    3.0  0.10 0.08164966
18 1e+01    3.0  0.16 0.09660918
19 1e+02    3.0  0.15 0.11785113
20 1e+03    3.0  0.18 0.13165612
21 1e-01    4.0  0.26 0.11737878
22 1e+00    4.0  0.10 0.08164966
23 1e+01    4.0  0.16 0.11737878
24 1e+02    4.0  0.16 0.11737878
25 1e+03    4.0  0.19 0.12866839

- best performance: 0.09

- Detailed performance results:
  cost gamma error dispersion
1 1e-01 0.5 0.20 0.14142136
2 1e+00 0.5 0.09 0.08755950
3 1e+01 0.5 0.10 0.08164966
4 1e+02 0.5 0.11 0.09944289
5 1e+03 0.5 0.14 0.13498971
6 1e-01 1.0 0.11 0.09944289
7 1e+00 1.0 0.10 0.08164966
8 1e+01 1.0 0.09 0.07378648
9 1e+02 1.0 0.14 0.12649111
```



```
table(true=dat[-train,"y"], pred=
predict(tune.out$best.model,newx=dat[-train,]))
```

	pred	
true	1	2
1	55	23
2	16	6

```
library(ISLR); dim(Khan$xtrain)
```

63 2308

```
dim(Khan$ytest)
```

20 2308

```
table(Khan$ytrain)
```

1	2	3	4
8	23	12	20

```
table(Khan$ytest)
```

1	2	3	4
3	6	6	5

```
dat=data.frame(x=Khan$xtrain,  
y=as.factor(Khan$ytrain))
```

```
out=svm(y~., data=dat, kernel="linear",cost=10)  
table(out$fitted, dat$y)
```

	1	2	3	4
1	8	0	0	0
2	0	23	0	0
3	0	0	12	0
4	0	0	0	20

```
dat.te=data.frame(x=Khan$ptest,  
y=as.factor(Khan$ytest))
```

```
pred.te=predict(out, newdata=dat.te)  
table(pred.te, dat.te$y)
```

pred.te	1	2	3	4
1	3	0	0	0
2	0	6	2	0
3	0	0	4	0
4	0	0	0	5